API Gateway Demo

Introduction

From the frontend application (developed specifically to test the resilience of my containers) you can login either as a **regular user** (without the privilege of viewing the health status of the microservices) or as an **administrator**.

The main (and only) page displays the general system status, retrieved from the health API, along with:

- A summary of the requests made by the currently logged-in user
- A testing section where you can choose which API and which microservice to invoke

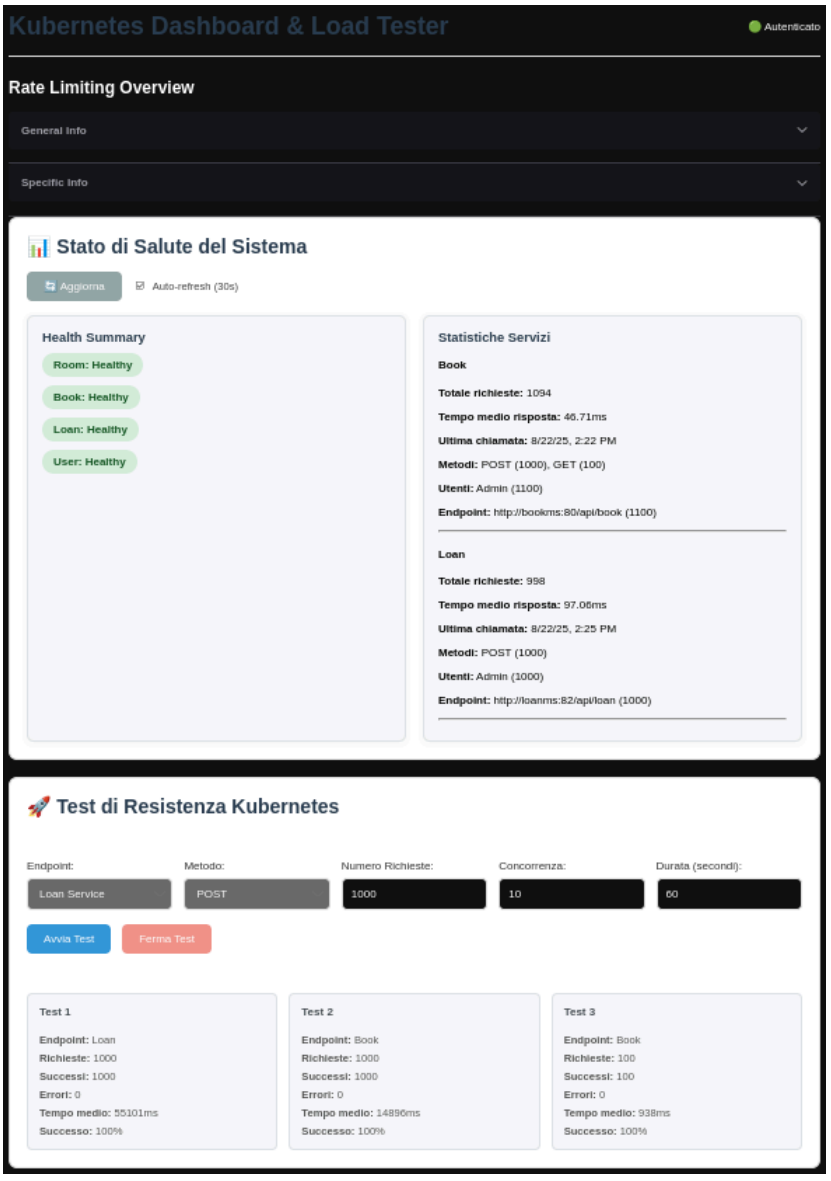

Usage example

After sending a high number of requests to the **Book** and **User** microservices:

We can observe that the **Horizontal Pod Autoscaler (HPA)** starts scaling the pods accordingly.

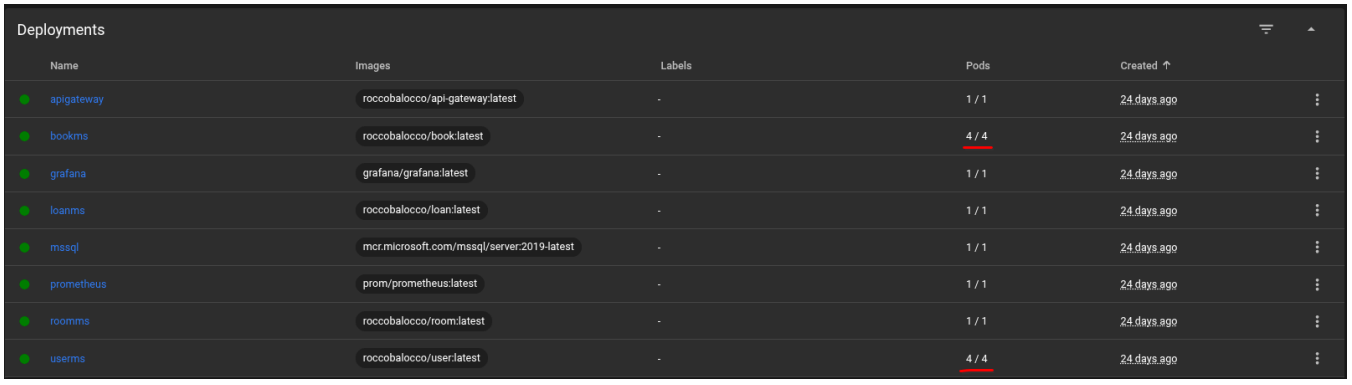

To stress-test the deployment, I sent **20,000 requests**, split between two different microservices.

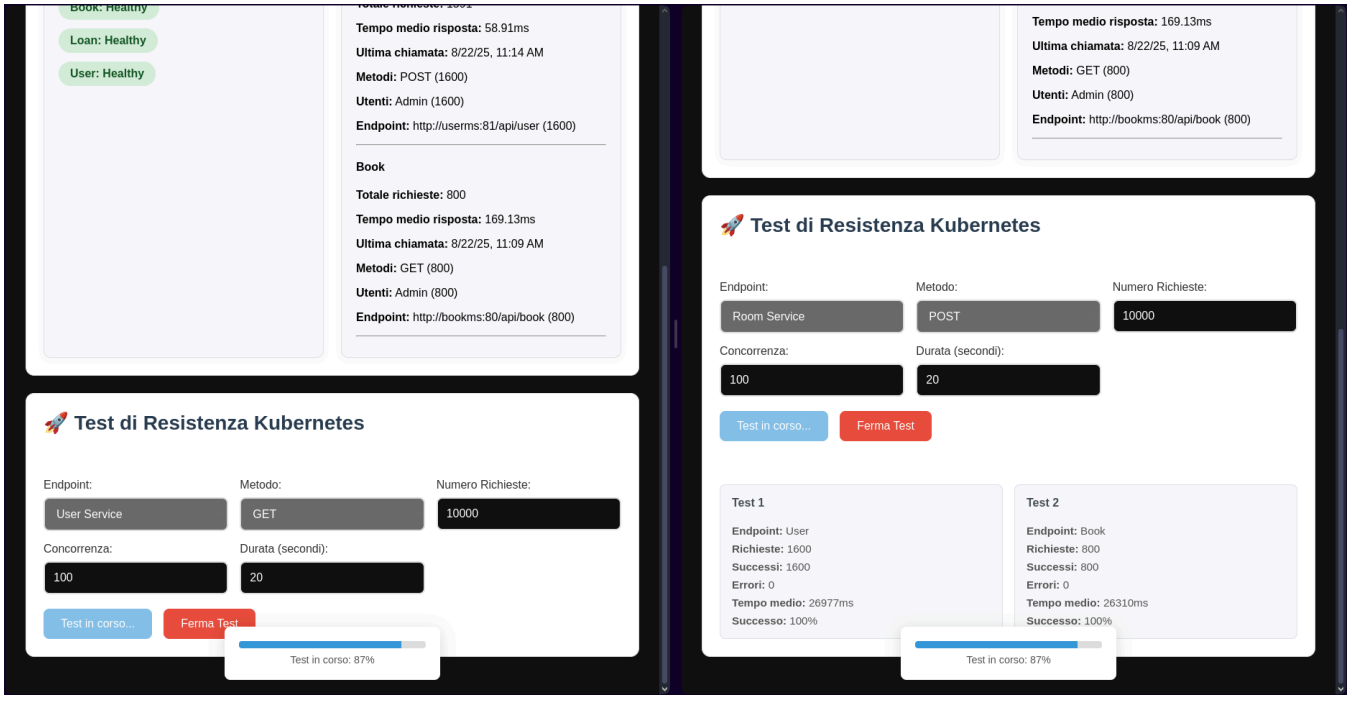

As a result, many pods crashed due to workload intensity and existing policy configurations.

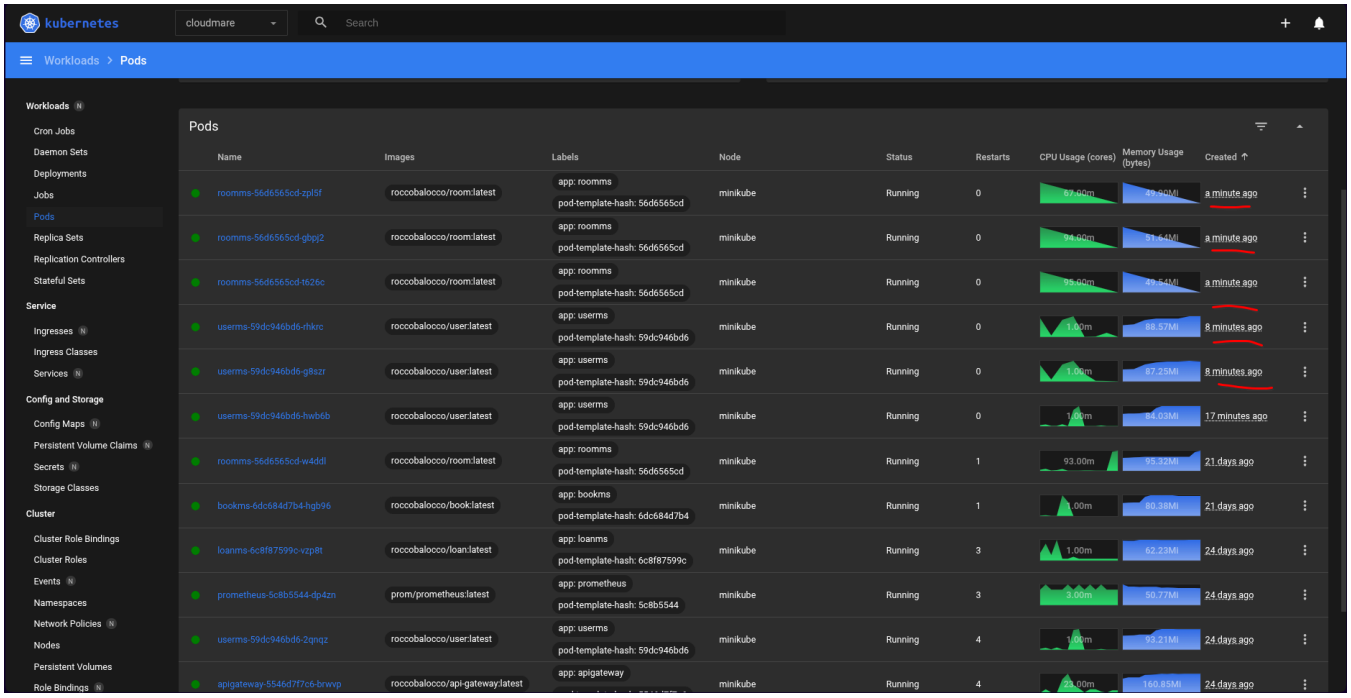

- The test on the **Room** service achieved a success rate of **17%**.

  Test 1

  Endpoint: Room
  Richieste: 10000
  Successi: 1741
  Errori: 8259
  Tempo medio: 16179ms
  Successo: 17%

- The test on the **User** service achieved a success rate of **14%**.

  Test 1

  Endpoint: User
  Richieste: 10000
  Successi: 1357
  Errori: 8643
  Tempo medio: 44320ms
  Successo: 14%

Although 10,000 requests were sent to each microservice, only a small number were processed successfully. The remaining requests were blocked due to circuit-breaking policies, pod crashes, or other constraints.

Statistiche Servizi

Room

Totale richieste: 1736
Tempo medio risposta: 43.17ms
Ultima chiamata: 8/22/25, 11:20 AM
Metodi: POST (1741)
Utenti: Admin (1741)
Endpoint: http://roomms:83/api/room (1741)

User

Totale richieste: 2944
Tempo medio risposta: 161.03ms
Ultima chiamata: 8/22/25, 11:23 AM
Metodi: POST (1600), GET (1357)
Utenti: Admin (2957)
Endpoint: http://userms:81/api/user (2957)

After this test, the system adjusted dynamically:

- Since the latest requests targeted the **User** microservice, the HPA scaled up its number of pods.
- At the same time, the number of pods for the **Room** microservice was scaled down.

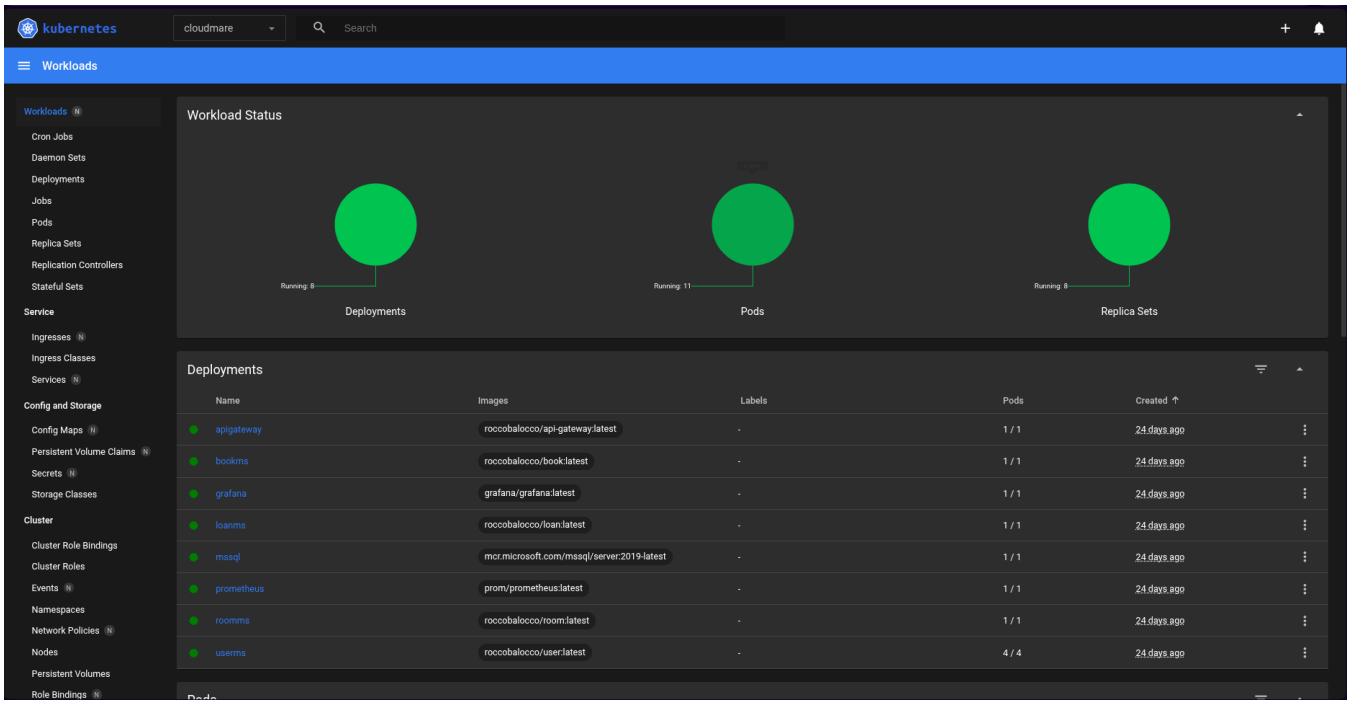

To trigger the **Per-IP** and **Global Fixed** policies in the Angular web application, you need to send many requests, as they are not sent simultaneously.

Prometheus

Thanks to the annotations included in the manifest files, **Prometheus** can scrape the metrics of the active pods.

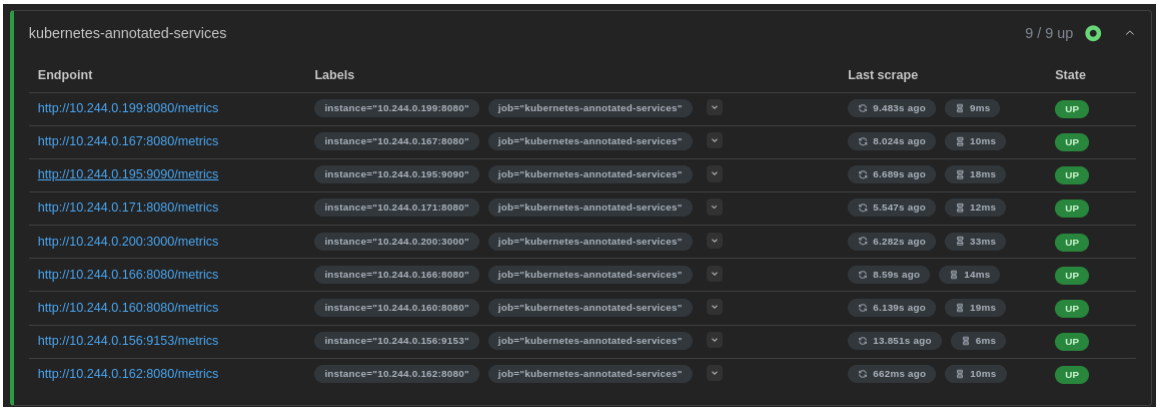

Grafana

Grafana connects to Prometheus as a data source and visualizes various performance and health metrics of the system.


Home › Drilldown › Metrics ›
Search...
ctrl+k
Data source
Prometheus
Filters
Filter by label values
Last 30 minutes
Quick search metrics
Sort by
Default
Grid
Rows
Group by labels
Search...
Selected: "endpoint"
clear
controller
datasource
datasource_type
db_name
deprecated_version
dialer_name
direction
dns_question_name
edition
endpoint
error_type
errorSource
event
event_name
event_severity
family
filter
frame_type
generation
goarch
goos
http_request_duration_seconds_bucket
Select
http_request_duration_seconds_count
sum(rate)
http_request_duration_seconds_sum
http_requests_in_progress
avg
http_requests_received_total
api/Gateway/book
(7/17)
14:15
14:20
14:25
14:30
14:35
14:40